

Tokyo Olympics 2020 Data Engineering Project (End-to-End)

Project Overview

This project demonstrates an end-to-end data engineering pipeline using Azure cloud services and Databricks to process and analyze Tokyo Olympics 2020 datasets. The primary objective was to extract raw Olympic data, transform and cleanse it, and load it into an Azure SQL Database for further analysis and reporting.

Tools and Technologies Used

- **Azure Data Lake Storage Gen2:** For storing raw datasets.
- **Azure Databricks:** For data transformation using PySpark.
- **Azure SQL Database:** Final destination for transformed data.
- **Azure Data Factory (ADF):** Used for orchestration and data ingestion (optional future enhancement).
- **PySpark:** Used for data manipulation and transformation.
- **JDBC:** For writing transformed data into Azure SQL Database.

Datasets Processed

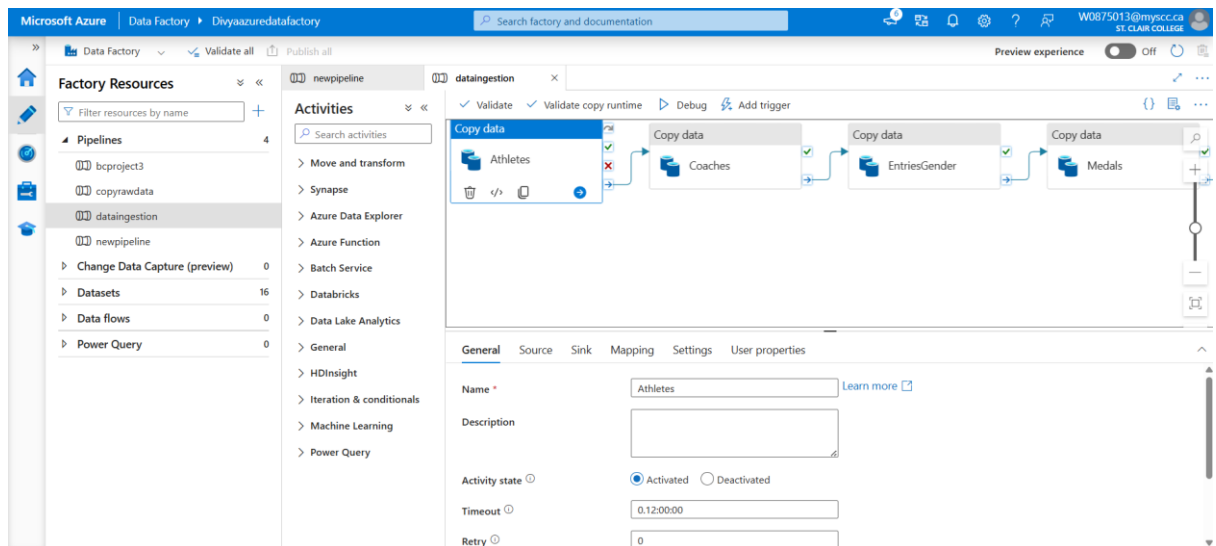
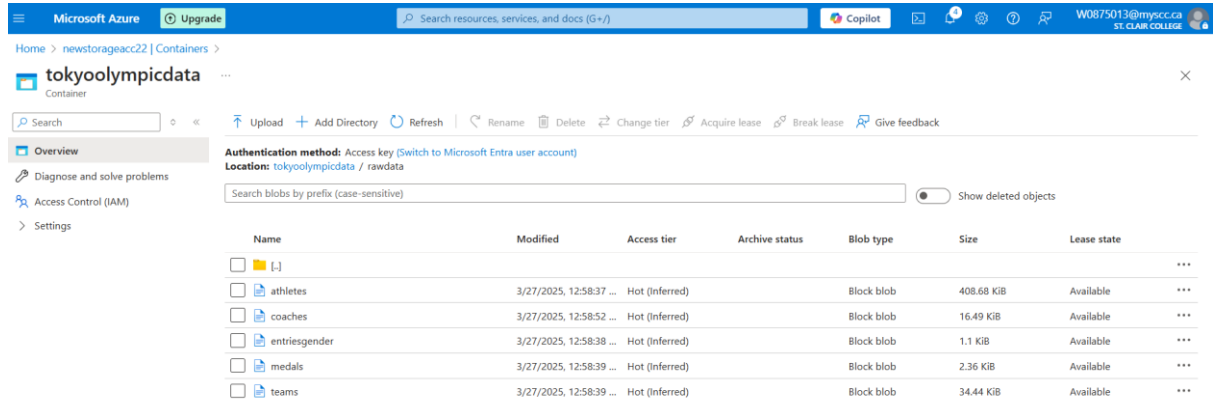
- Athletes
- Teams
- Entries by Gender
- Coaches
- Medals

Steps Implemented

1. Data Ingestion

- Uploaded Tokyo Olympics datasets (CSV files) to Azure Data Lake Storage Gen2.
- Created a bronze layer folder for raw files.

- Created a pipeline using copy activity to transfer files to transformeddata folder.



2. Data Transformation with Databricks

- Read raw CSVs using PySpark into DataFrames.
- Applied necessary transformations like schema definition, null handling, and finding the top 10 countries with the highest number of gold medals
- Finding the total number of athletes per country
- Validated each DataFrame to ensure data quality and schema consistency.

```

▶ ✓ Mar 27, 2025 (24s) 1

# Define variables
storage_account_name = "newstorageacc22"
container_name = "tokyoolympicdata"
directory_name = "rawdata"

# Storage account access key
access_key = "bxwas0MbIT5D+YfV/RcEIcThfWLZSTnRTtPUUvDBkqRBd8JPgRNF9mqjWV9GXV97vscL1v9pT7tv+ASthFeLoA=="

# Mount path
mount_point = f"/mnt/{container_name}"

# Mount ADLS to Databricks
dbutils.fs.mount(
    source=f"wasbs://{container_name}@{storage_account_name}.blob.core.windows.net",
    mount_point=mount_point,
    extra_configs={f"fs.azure.account.key.{storage_account_name}.blob.core.windows.net": access_key}
)

# Verify mount
display(dbutils.fs.ls(mount_point))

▶ (2) Spark Jobs

```

Table +

🔍 🔍 📄 📄

```

▶ ✓ 10:52 PM (16s) 3

# Read the "athletes" file
df_athletes = spark.read.csv("dbfs:/mnt/tokyoolympicdata/rawdata/athletes", header=True, inferSchema=True)
df_athletes.show(5)

▶ (3) Spark Jobs

```

df_athletes: pyspark.sql.dataframe.DataFrame = [PersonName: string, Country: string ... 1 more field]

PersonName	Country	Discipline
AALERUD Katrine	Norway	Cycling Road
ABAD Nestor	Spain	Artistic Gymnastics
ABAGNALE Giovanni	Italy	Rowing
ABALDE Alberto	Spain	Basketball
ABALDE Tamara	Spain	Basketball

only showing top 5 rows

```

▶ ✓ 10:58 PM (1s) 5

# Read the entriessgender CSV file
df_entriessgender = spark.read.csv("dbfs:/mnt/tokyoolympicdata/rawdata/entriessgender", header=True, inferSchema=True)

# Print the schema
df_entriessgender.printSchema()

▶ (2) Spark Jobs

```

df_entriessgender: pyspark.sql.dataframe.DataFrame = [Discipline: string, Female: integer ... 2 more fields]

root

```

|-- Discipline: string (nullable = true)
|-- Female: integer (nullable = true)
|-- Male: integer (nullable = true)
|-- Total: integer (nullable = true)

```

```
10:58 PM (1s) 7

# Finding the top 10 countries with the highest number of gold medals
from pyspark.sql.functions import col

# Sort by GoldMedals in descending order
df_top_gold = df_medals.select("TeamCountry", "Gold").orderBy(col("Gold").desc())

# Show top 10 countries
df_top_gold.show(10)
```

▶ (1) Spark Jobs

▶ df_top_gold: pyspark.sql.dataframe.DataFrame = [TeamCountry: string, Gold: integer]

TeamCountry	Gold
United States of ...	39
People's Republic...	38
Japan	27
Great Britain	22
ROC	20
Australia	17
Netherlands	10
France	10
Germany	10
Italy	10

only showing top 10 rows

```
10:58 PM (4s) 8 Python

# Find the Total Number of Athletes per Country
df_athletes = spark.read.csv("dbfs:/mnt/tokyoolympicdata/rawdata/athletes", header=True, inferSchema=True)

from pyspark.sql.functions import count

df_total_athletes = df_athletes.groupBy("Country").agg(count("*").alias("Total_Athletes"))
df_total_athletes.show(10)
```

▶ (4) Spark Jobs

▶ df_athletes: pyspark.sql.dataframe.DataFrame = [PersonName: string, Country: string ... 1 more field]

▶ df_total_athletes: pyspark.sql.dataframe.DataFrame = [Country: string, Total_Athletes: long]

Country	Total_Athletes
Chad	3
Paraguay	8
Yemen	3
Islamic Republic ...	66
Chinese Taipei	67
Senegal	9
Sweden	129
Kiribati	3
Republic of Korea	223
Guyana	7

3. JDBC Connection Setup

```
jdbc_url =
"jdbc:sqlserver://divyaproject.database.windows.net:1433;database=bcproject3;encrypt=true;tr
ustServerCertificate=false;hostNameInCertificate=*.database.windows.net;loginTimeout=30
;"
```

```
connection_properties = {  
    "user": "Yourusername",  
    "password": "Your password",  
    "driver": "com.microsoft.sqlserver.jdbc.SQLServerDriver"  
}
```

4. Data Loading to Azure SQL Database

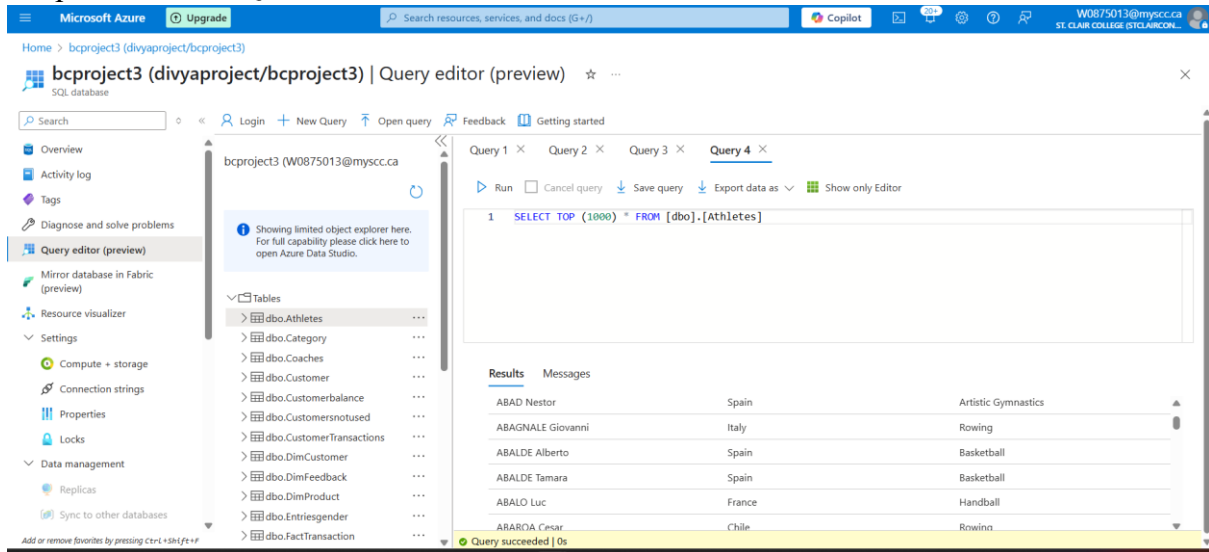
Transformed DataFrames were written to Azure SQL Database using the `.write.jdbc()` method:

```
df_athletes.write.jdbc(url=jdbc_url, table="Athletes", mode="overwrite",  
properties=connection_properties)  
  
df_teams.write.jdbc(url=jdbc_url, table="Teams", mode="overwrite",  
properties=connection_properties)  
  
df_entriesgender.write.jdbc(url=jdbc_url, table="EntriesGender", mode="overwrite",  
properties=connection_properties)  
  
df_coaches.write.jdbc(url=jdbc_url, table="Coaches", mode="overwrite",  
properties=connection_properties)  
  
df_medals.write.jdbc(url=jdbc_url, table="Medals", mode="overwrite",  
properties=connection_properties)
```

```
df_medals.write.jdbc(  
    url=jdbc_url,  
    table="Medals",  
    mode="overwrite",  
    properties=connection_properties  
)
```

► (1) Spark Jobs

Output from the SQL Database for Athletes file:



The screenshot shows the Microsoft Azure portal interface for a SQL database named 'bcproject3'. The 'Query editor (preview)' is open, displaying a query that selects the top 1000 records from the 'dbo.Athletes' table. The query is: `1 SELECT TOP (1000) * FROM [dbo].[Athletes]`. The results are displayed in a table with columns for Name, Country, and Sport. The first five rows are: ABAD Nestor (Spain, Artistic Gymnastics), ABAGNALE Giovanni (Italy, Rowing), ABALDE Alberto (Spain, Basketball), ABALDE Tamara (Spain, Basketball), and ABALO Luc (France, Handball). The status bar at the bottom indicates 'Query succeeded | 0s'.

Name	Country	Sport
ABAD Nestor	Spain	Artistic Gymnastics
ABAGNALE Giovanni	Italy	Rowing
ABALDE Alberto	Spain	Basketball
ABALDE Tamara	Spain	Basketball
ABALO Luc	France	Handball

- Similarly, other tables like medals, teams, coaches and entriesgender were created.

Outcome

- Successfully built and tested an ETL pipeline that processes and stores Tokyo Olympics data into an Azure SQL Database.
- Created a scalable and repeatable process that can be extended to other Olympic datasets.